

Reusable backend resident performance

18,141 targets $\times$ 3 experiments $\cdot$ n = 5 requests per standalone arm

Python-fast / bundle load: 4.26 $\times$

